Machine Intelligence 1 **Exercise Sheet 4**
WS 2019/20, Obermayer/Kashef due: 14.11.2019 at 23:55

Gradient methods for parameter optimization

Exercise T4.1: Multilayer perceptron (tutorial)

(a) Recap the optimization of the MLP parameters (via the backpropagation algorithm).

(b) Outline the weight space symmetries giving rise to $\Pi_{v=1}^{L} N_v! \cdot 2^{N_v}$ equivalent solutions where L is the number of hidden layers and N_v the respective number of neurons in layer v $\Longrightarrow$ no unique global minimum but a large equivalence class of (best) solutions.

Exercise T4.2: Linear neuron for regression (tutorial)

To prepare for the homework, we discuss a simple connectionist neuron with linear output function for a real one-dimensional input $x \in \mathbb{R}$ and output $y \in \mathbb{R}$.

(a) Describe the output function $y(x; \underline{\mathbf{w}})$ of the neuron in vector notation.

(b) Derive the gradient and Hessian matrix of the quadratic error function.

(c) Solve the optimization of the quadratic error function for a data set $\{(x^{(\alpha)}, y_T^{(\alpha)})\}_{\alpha=1,\ldots,p}$ analytically in matrix form.

(d) Calculate the solution when the objective includes the quadratic training cost E^T plus a "weight decay" regularization term as used in *ridge regression*, i.e.

$$R_{[\underline{\mathbf{w}}]} = E^T_{[\underline{\mathbf{w}}]} + \lambda \, ||\underline{\mathbf{w}}||^2$$

Exercise T4.3: Conjugate gradient (tutorial)

(a) How does the convergence speed of *gradient descent* depend on the learning rate η?

(b) Describe how *line search* speeds up convergence.

(c) What is a *conjugate direction* and how can it improve convergence speed?

Exercise H4.1: Line search (homework, 4 points)

In this exercise you will analyze line search based on the simple example of a linear neuron with quadratic cost function $E^T_{[\underline{\mathbf{w}}]}$. Here we optimize the cost function along a given direction $\underline{\mathbf{d}}_t$ (that can be but is not necessarily identical to the gradient):

$$\underline{\mathbf{w}}_{t+1} = \underline{\mathbf{w}}_t - \eta_t \, \underline{\mathbf{d}}_t \, .$$

(a) (1 point) Derive the 2nd order Taylor approximation of an *arbitrary* $E^T_{[\underline{\mathbf{w}}_{t+1}]}$ around $\underline{\mathbf{w}}_t$.

(b) (1 point) Derive a bound on the step size η_t using the above approximation in $E^T_{[\underline{\mathbf{w}}_{t+1}]} \overset{!}{\leq} E^T_{[\underline{\mathbf{w}}_t]}$.

(c) (1 point) Derive the optimal step size η_t^* for the quadratic cost function

$$E^T_{[\underline{\mathbf{w}}]} = \tfrac{1}{2}(\underline{\mathbf{w}} - \underline{\mathbf{w}}^*)^\top \underline{\mathbf{H}}(\underline{\mathbf{w}} - \underline{\mathbf{w}}^*)$$

by minimizing the cost function w.r.t. η. Make sure your solution depends only on known quantities like the weight vector $\underline{\mathbf{w}}_t$, the gradient $\underline{\nabla} E^T_{[\underline{\mathbf{w}}_t]}$ and/or the Hessian $\underline{\mathbf{H}}$ of $E^T_{[\underline{\mathbf{w}}_t]}$.

(d) (1 point) Prove that the gradient $\underline{\nabla} E^T_{[\underline{\mathbf{w}}_{t+1}]}$ after one update step with *line search* is orthogonal to the optimized direction $\underline{\mathbf{d}}_t$.

Solution

(a) Let the gradient $\underline{\mathbf{g}}_t := \underline{\nabla} E^T_{[\underline{\mathbf{w}}_t]}$ and the Hessian matrix $\underline{\mathbf{H}}_t := \underline{\boldsymbol{\Delta}} E^T_{[\underline{\mathbf{w}}_t]}$, then

$$\begin{aligned} E^T_{[\underline{\mathbf{w}}_{t+1}]} &\approx E^T_{[\underline{\mathbf{w}}_t]} + (\underline{\mathbf{w}}_{t+1} - \underline{\mathbf{w}}_t)^\top \underline{\mathbf{g}}_t + \tfrac{1}{2}(\underline{\mathbf{w}}_{t+1} - \underline{\mathbf{w}}_t)^\top \underline{\mathbf{H}}_t (\underline{\mathbf{w}}_{t+1} - \underline{\mathbf{w}}_t) \\ &= E^T_{[\underline{\mathbf{w}}_t]} - \eta_t \, \underline{\mathbf{d}}_t^\top \underline{\mathbf{g}}_t + \frac{\eta^2}{2} \underline{\mathbf{d}}_t^\top \underline{\mathbf{H}}_t \underline{\mathbf{d}}_t \,. \end{aligned}$$

(b) $E^T_{[\underline{\mathbf{w}}_{t+1}]} \approx E^T_{[\underline{\mathbf{w}}_t]} - \eta_t \, \underline{\mathbf{d}}_t^\top \underline{\mathbf{g}}_t + \frac{\eta_t^2}{2} \underline{\mathbf{d}}_t^\top \underline{\mathbf{H}}_t \underline{\mathbf{d}}_t \overset{!}{\leq} E^T_{[\underline{\mathbf{w}}_t]} \quad \Rightarrow \quad \eta_t \overset{!}{\leq} 2 \frac{\underline{\mathbf{d}}_t^\top \underline{\mathbf{g}}_t}{\underline{\mathbf{d}}_t^\top \underline{\mathbf{H}}_t \underline{\mathbf{d}}_t} \,.$

(c) Solve $\min_\eta E^T_{(\underline{\mathbf{w}}_t + 1)}$ by setting the derivative w.r.t. η to zero:

$$\begin{aligned} \frac{\partial E^T_{[\underline{\mathbf{w}}_{t+1}]}}{\partial \eta} &= \left(\frac{\partial E^T_{[\underline{\mathbf{w}}_{t+1}]}}{\partial \underline{\mathbf{w}}_{t+1}} \right)^\top \frac{\partial \underline{\mathbf{w}}_{t+1}}{\partial \eta} = (\underline{\mathbf{H}}_t \underline{\mathbf{w}}_t - \eta_t \underline{\mathbf{H}}_t \underline{\mathbf{d}}_t - \underline{\mathbf{H}}_t \underline{\mathbf{w}}^*)^\top (-\underline{\mathbf{d}}_t) \overset{!}{=} 0 \\ \Rightarrow \quad \eta^* &= \frac{\underline{\mathbf{d}}_t^\top \underline{\mathbf{H}}_t (\underline{\mathbf{w}}_t - \underline{\mathbf{w}}^*)}{\underline{\mathbf{d}}_t^\top \underline{\mathbf{H}}_t \underline{\mathbf{d}}_t} = \frac{\underline{\mathbf{d}}_t^\top \underline{\mathbf{g}}_t}{\underline{\mathbf{d}}_t^\top \underline{\mathbf{H}}_t \underline{\mathbf{d}}_t} \,, \qquad \text{as} \quad \underline{\mathbf{g}}_t = \underline{\mathbf{H}}_t (\underline{\mathbf{w}}_t - \underline{\mathbf{w}}^*) \,. \end{aligned}$$

(d) The gradient is orthogonal to the direction if $\underline{\mathbf{d}}_t^\top \underline{\mathbf{g}}_{t+1} = 0$.

$$\begin{aligned} \underline{\mathbf{g}}_{t+1} &= \underline{\mathbf{H}}_t (\underline{\mathbf{w}}_{t+1} - \underline{\mathbf{w}}^*) = \underline{\mathbf{H}}_t (\underline{\mathbf{w}}_t - \eta_t \underline{\mathbf{d}}_t - \underline{\mathbf{w}}^*) = \underline{\mathbf{g}}_t - \eta_t \underline{\mathbf{H}}_t \underline{\mathbf{d}}_t \\ \underline{\mathbf{d}}_t^\top \underline{\mathbf{g}}_{t+1} &= \underline{\mathbf{d}}_t^\top \underline{\mathbf{g}}_t - \frac{\underline{\mathbf{d}}_t^\top \underline{\mathbf{g}}_t}{\underline{\mathbf{d}}_t^\top \underline{\mathbf{H}}_t \underline{\mathbf{d}}_t} \underline{\mathbf{d}}_t^\top \underline{\mathbf{H}}_t \underline{\mathbf{d}}_t = 0 \end{aligned}$$

Exercise H4.2: Comparison of gradient descent methods (homework, 6 points)

In this exercise we compare the performance of three learning procedures applied to a simple connectionist neuron with a linear output function. All procedures will compute the gradient using the entire training set (batch gradient descent). The procedures are: (i) Gradient (or steepest) descent with constant learning rate, (ii) steepest descent combined with a line search method to determine the learning rate, and (iii) the conjugate gradient method.

Training Data: The training data set consists of three points ($p = 3$):

$$\{(x^{(\alpha)}, y_T^{(\alpha)})\} = \{(-1, -0.1), (0.3, 0.5), (2, 0.5)\},$$

i.e. for a given data point, both input and output are scalar values.

Cost function: The gradient for the *quadratic error* function is given by

$$\underline{\mathbf{g}} = \frac{\partial E^T}{\partial \underline{\mathbf{w}}} = \underline{\mathbf{H}}\,\underline{\mathbf{w}} + \underline{\mathbf{b}}, \qquad \text{with} \quad \underline{\mathbf{H}} = \underline{\mathbf{X}}\,\underline{\mathbf{X}}^\top \quad \text{and} \quad \underline{\mathbf{b}} = -\underline{\mathbf{X}}\,\underline{\mathbf{y}}^\top,$$

where $\underline{\mathbf{X}} = \begin{pmatrix} 1 & 1 & \dots & 1 \\ x^{(1)} & x^{(2)} & \dots & x^{(p)} \end{pmatrix} \in \mathbb{R}^{2,p}$ and $\underline{\mathbf{y}} = \left(y_T^{(1)}, y_T^{(2)}, \dots, y_T^{(p)}\right) \in \mathbb{R}^{1,p}$.

Initialization: Use the following initialization for all three (batch) gradient methods:

$$\underline{\mathbf{w}}_1 = (w_0, w_1)_1^\top = (-0.45, 0.2)^\top$$

(a) (2 points) *Gradient Descent:* Implement a steepest descent procedure where the weights at iteration $t + 1$ are calculated using the weights and the gradient at iteration t

$$\underline{\mathbf{w}}_{t+1} = \underline{\mathbf{w}}_t - \eta\,\underline{\mathbf{g}}_t,$$

with an adequate learning rate η and where $\underline{\mathbf{g}}_t = \underline{\mathbf{g}}(\underline{\mathbf{w}}_t)$. Plot the resulting weight vectors from all iterations as a scatter plot (w_0 vs. w_1), and in an additional plot (w_i vs. iterations t), to show the development of each parameter during gradient descent.

(b) (2 points) *Line Search:* Implement a line search procedure

$$\underline{\mathbf{w}}_{t+1} = \underline{\mathbf{w}}_t - \eta\,\underline{\mathbf{g}}_t, \qquad \text{with optimal step size} \qquad \eta = \frac{\underline{\mathbf{g}}_t^\top \underline{\mathbf{g}}_t}{\underline{\mathbf{g}}_t^\top \underline{\mathbf{H}} \underline{\mathbf{g}}_t}.$$

Plot the resulting weight vectors from all iterations as a scatter plot (w_0 vs. w_1), and in an additional plot (w_i vs. iterations t), to show the development of the parameters during line search.

(c) (2 points) *Conjugate Gradient:* Implement a conjugate gradient procedure:
Initialize: $\underline{\mathbf{w}}_1, \underline{\mathbf{d}}_1 = -\underline{\mathbf{g}}_1$
while *stopping criterion not satisfied* **do**
  minimize E along $\underline{\mathbf{d}}_t$: $\underline{\mathbf{w}}_{t+1} = \underline{\mathbf{w}}_t + \eta_t \underline{\mathbf{d}}_t$ with step size $\eta_t = -\frac{\underline{\mathbf{d}}_t^\top \underline{\mathbf{g}}_t}{\underline{\mathbf{d}}_t^\top \underline{\mathbf{H}}\,\underline{\mathbf{d}}_t}$
  calculate new gradient $\underline{\mathbf{g}}_{t+1} = \underline{\mathbf{H}}\,\underline{\mathbf{w}}_{t+1} + \underline{\mathbf{b}}$
  calculate new conjugate direction $\underline{\mathbf{d}}_{t+1} = \underline{\mathbf{g}}_{t+1} + \beta_t \underline{\mathbf{d}}_t$ with "momentum"

$$\beta_t = -\frac{\underline{\mathbf{g}}_{t+1}^\top \underline{\mathbf{g}}_{t+1}}{\underline{\mathbf{g}}_t^\top \underline{\mathbf{g}}_t}. \qquad \text{(Fletcher-Reeves form)}$$

  increase $t \leftarrow t + 1$
end

Plot the resulting weight vectors from all iterations as a scatter plot (w_0 vs. w_1), and in an additional plot (w_i vs. iterations t), to show the development of the parameters during conjugate gradient descent.

Compare the different methods in terms of convergence behaviour.

Solution

(a) Plot of the weight trace and their evolution over time – fairly slow convergence.

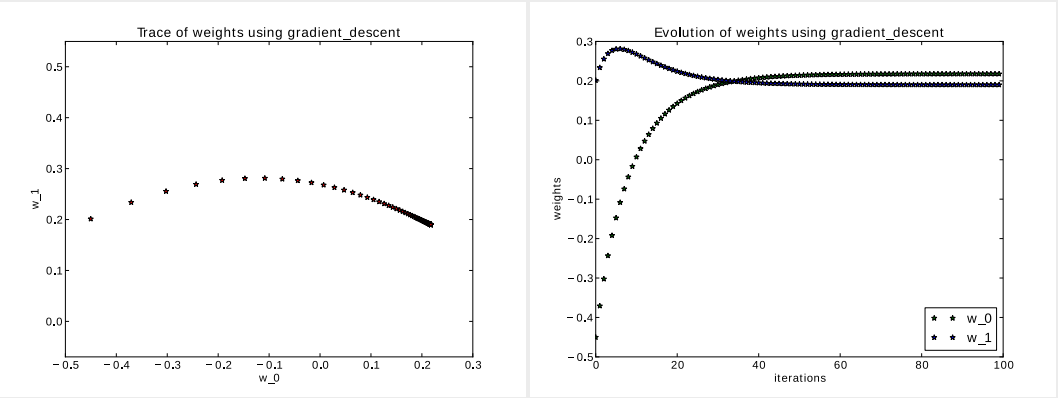

(b) Plot of the weight trace and their evolution over time – much faster convergence in the beginning, noticeable slower close to the minimum.

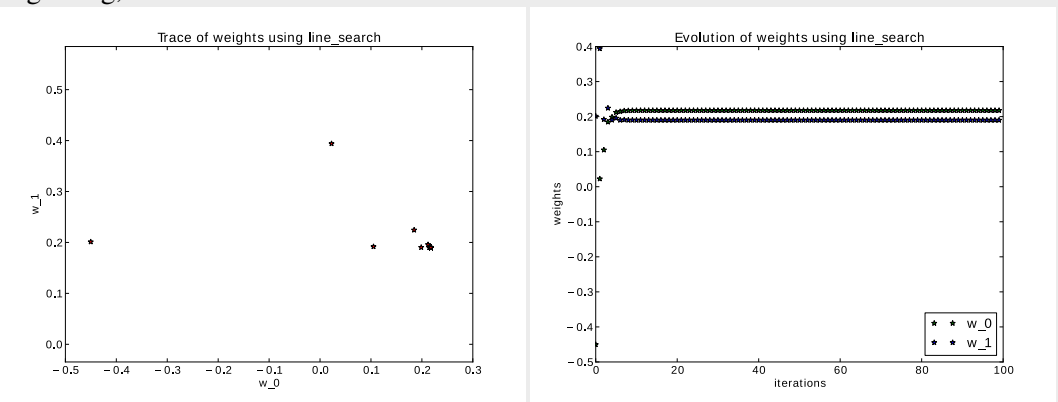

(c) Plot of the weight trace and their evolution over time – convergence in 2 iterations.

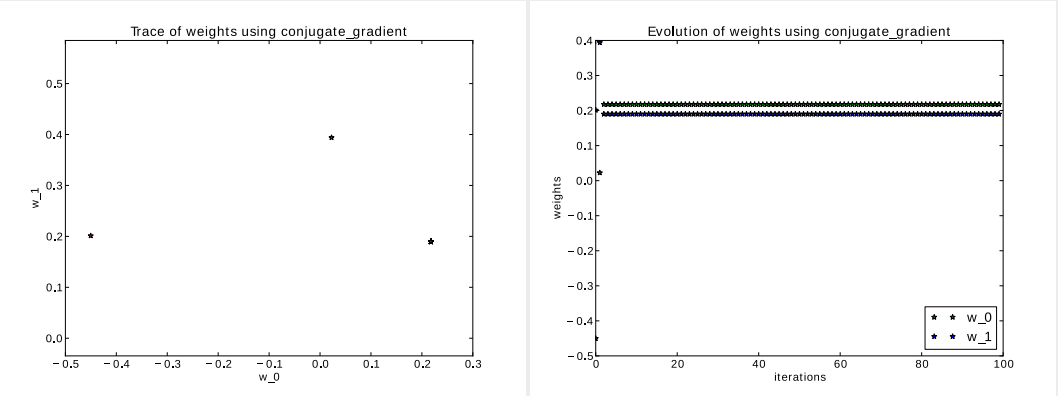

Total 10 points.